tqdm: three independent nanolize runs

Saved-candidate trajectory: source + configuration

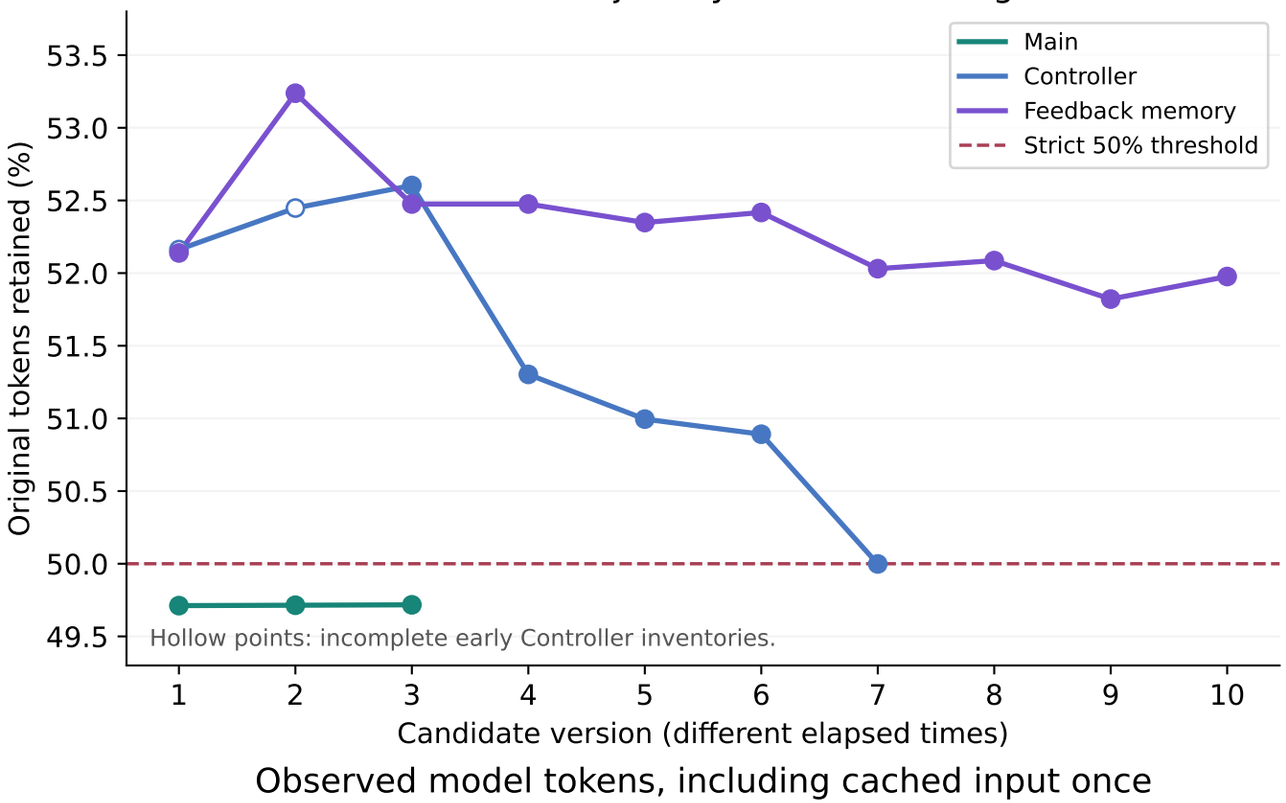

Final size: all three exceed 50% source retention

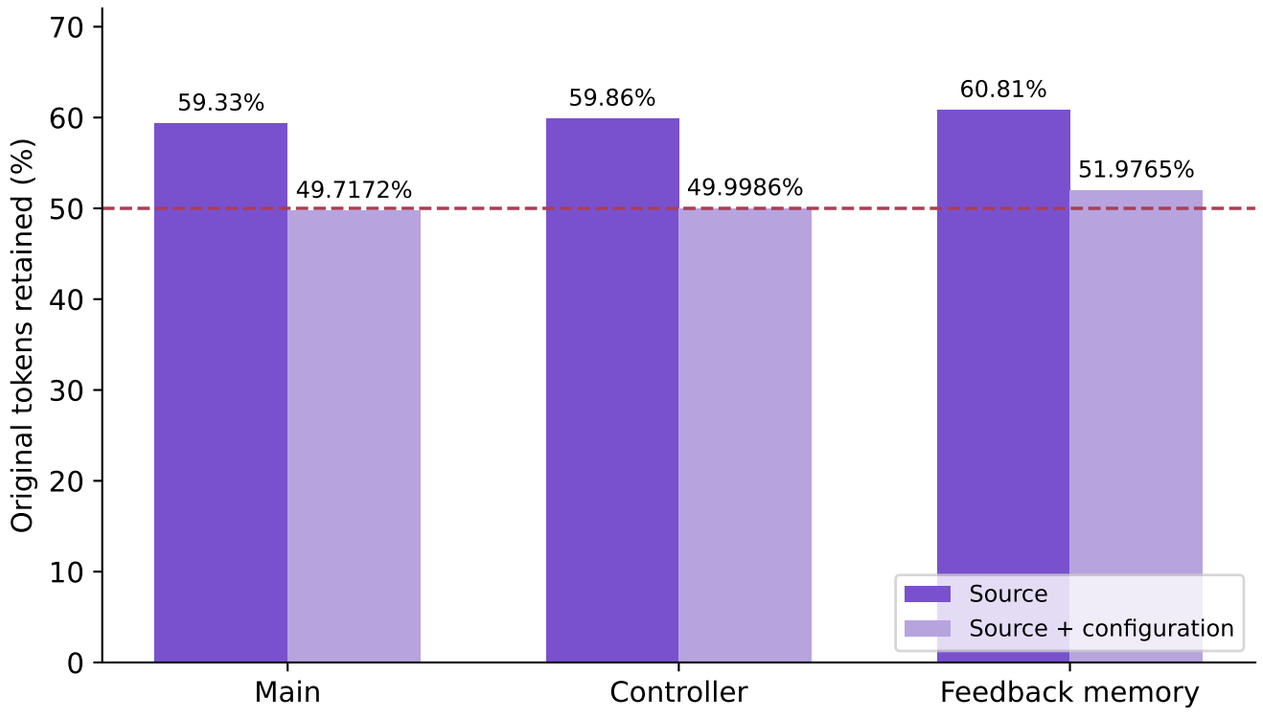

Observed model tokens, including cached input once

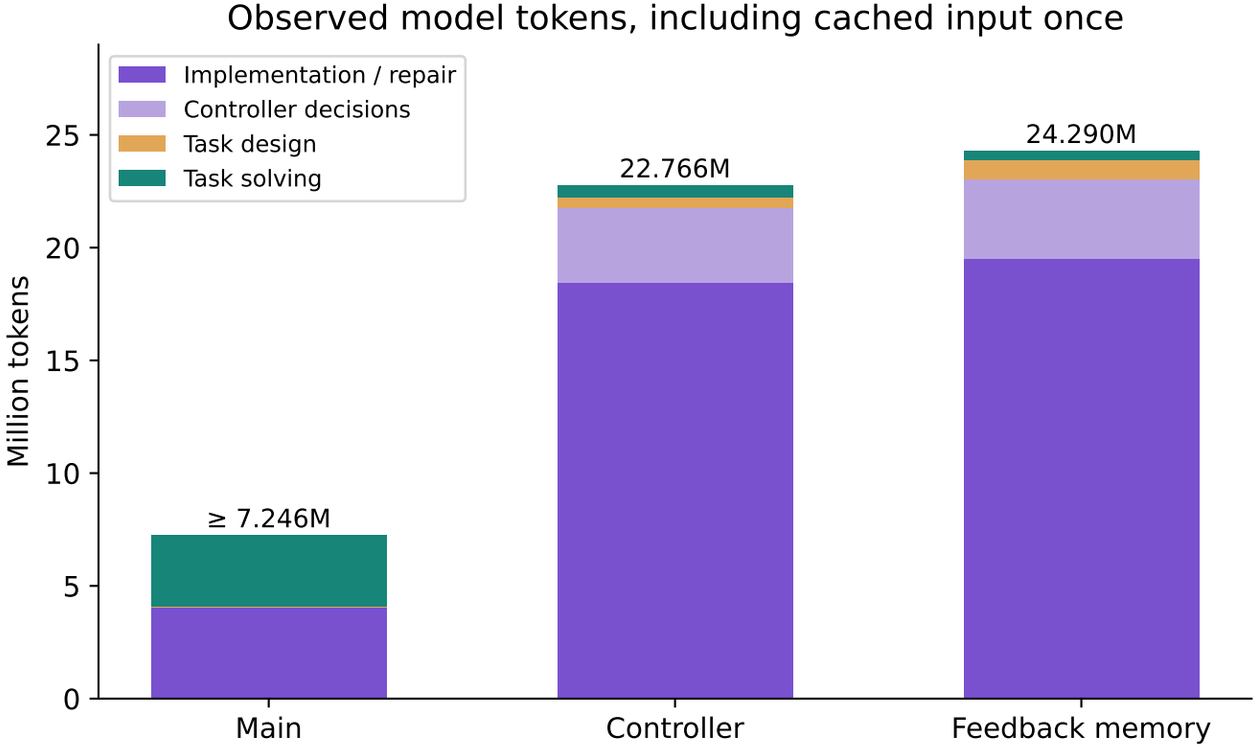

Wall time: submission to final harness outcome

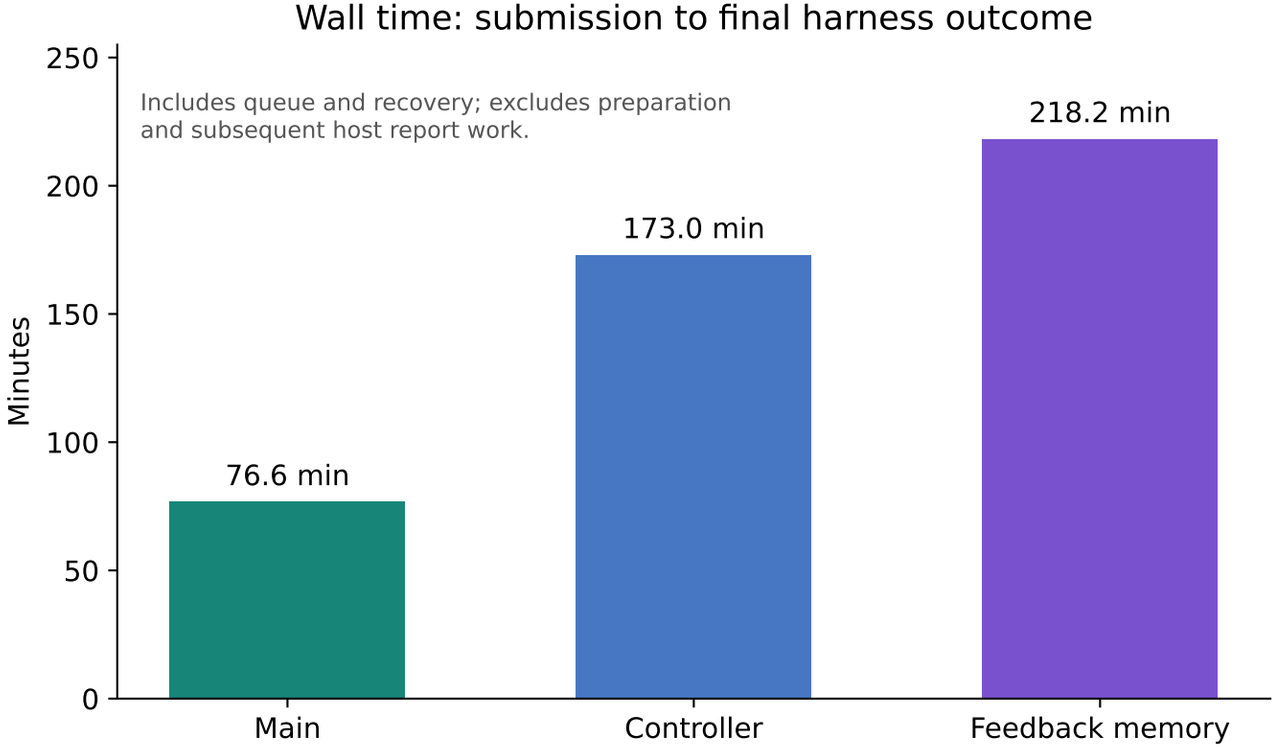

Different gates, task budgets and owner interventions: observational comparison, not a controlled ablation.
Main has one unmetered interrupted author call; about 91–92% of observed tokens are cached input.